Ananya Rachuri

rachuriananya5@gmail.com | 4698912868 | LinkedIn | GitHub

Senior Data Analyst with 5+ years of experience transforming complex, multi-source datasets into actionable insights that drive decision-making across global, matrixed organizations. Proven expertise in data governance, DQC frameworks, audit-ready compliance reporting, and automation including co-building an internal tool that reduced manual operations by 60%. Skilled in Power BI, Tableau, SQL, Python, and ETL pipeline development, with a track record of owning end-to-end analytics solutions for cross-functional business stakeholders.

EDUCATION

Ph.D., Artificial Intelligence | Starting Fall 2026 *School TBD*
Research: AI-Driven Data Quality Automation & Intelligent Analytics Systems

The University of Texas at Dallas | ***Master's in Information Technology Management***
Coursework: Database Systems, Business Intelligence, Machine Learning, Data Governance

SKILLS

Languages & Query: Python (PySpark, Pandas, NumPy), SQL (Advanced), R, Spark SQL, DAX, FastAPI, REST APIs, Git/GitHub.
Analytics & Visualization: Power BI, Tableau, Looker, Google Data Studio, Excel (Power Query, VBA, Pivot Tables).
Cloud & Data Platforms: Azure (ADF, Synapse, DevOps), AWS (S3, Redshift, Lambda), GCP (BigQuery, Composer), Snowflake, Databricks, Delta Lake.
AI / ML: scikit-learn, XGBoost, Logistic Regression, LLM integration (Claude API), Streamlit, Generative AI, anomaly detection.
Data Engineering & ETL: Apache Airflow, Apache Kafka, SSIS, ETL/ELT pipelines, CI/CD (GitHub Actions, Azure DevOps), Docker, Kubernetes.
Data Governance & Quality: MDM/PMDM (Master Data Management), DQC (Data Quality Controls), data lineage, schema enforcement, data profiling, GAP analysis, audit readiness, data migration, data cleansing, SOX/FFIEC.
Business & Delivery: Agile/Scrum (JIRA, Confluence), stakeholder management, data storytelling, executive reporting, cross-functional collaboration, global/matrixed organization, actionable insights, decision-making support.

WORK EXPERIENCE

TekMind Systems LLC McKinney, TX
Senior Data Engineer *Jun 2025 – Present*

- Architected PySpark and Spark SQL feature pipelines in Databricks supporting 1 ML team across 3 concurrent projects, cutting data science prep time by ~35% and accelerating model iteration cycles by ~20% — transforming raw data into actionable insights that drive decision-making.
- Designed Delta Lake tables with schema enforcement, optimized partitioning, and ACID-compliant writes; built CI/CD pipelines via GitHub Actions reducing release cycle time by ~40% across Dev/QA/Prod environments, recommending and implementing improvements to data acquisition and system efficiency.
- Enforced 99.8% SLA adherence across end-to-end pipeline schedules using Databricks Jobs and Workflows, with proactive alerting for data freshness violations before downstream ML consumption.
- Co-delivered sprint review and project status reporting to cross-functional stakeholders, translating pipeline health and data quality metrics into business-readable progress updates; identified patterns and trends in complex datasets to support operational excellence.

Cigna Healthcare Hartford, CT (Remote)
Market Insights Analyst *Jul 2024 – May 2025*

- Co-built "Doctor," an internal automation tool for pipeline monitoring and alerting, reducing manual operations by 60% across VOC ETL pipelines (Qualtrics, Clarabridge, Verint) processing 2M+ records/month; led automation project end-to-end in collaboration with the DTC team.

- Independently flagged and resolved data quality issues through data cleansing, readiness tracking, and validation frameworks, achieving 95%+ accuracy across analytics-ready datasets and delivering actionable insights that contributed to a 12% improvement in customer retention forecast accuracy.
- Supported audit-related data preparation and FFIEC compliance reporting with zero audit failures across 11 months, partnering with the compliance owner to standardize data content, validate datasets, and deliver audit-ready outputs on schedule.
- Prepared executive dashboard data and quality-validated reporting inputs for senior stakeholders across a global, matrixed organization, supporting data-driven decision-making through clean, analysis-ready data delivery.

Avis Solutions San Francisco, CA
Data Analyst *Jan 2022 – Dec 2023*

- Owned end-to-end design, maintenance, and iteration of Power BI and Tableau dashboards for Shipium, gathering requirements directly from business stakeholders and cutting ad-hoc report requests by 30% through self-service analytics adoption across 500K+ users.
- Built and maintained Spark SQL pipelines with optimized partitioned tables, improving batch SLA reliability by 22% and ensuring consistent data freshness for downstream dashboard consumption.
- Designed data ingestion, harmonization, and data migration workflows integrating 8 multi-source systems (AWS S3, Redshift, Azure Blob); developed and maintained Data Quality Controls (DQC) achieving 99.1% data accuracy with clean audit trails across all integrated sources.
- Translated business requirements gathered directly from analysts, architects, and product owners into scalable Databricks pipeline solutions across Agile sprints; collaborated with cross-functional teams and Data Domain partners to address business needs and standardize data processes.

Mindtree India
Junior Data Engineer *Jan 2021 – Dec 2021*

- Developed Python and SQL automation workflows across AWS, GCP, and Azure, executing GAP analysis on 15+ assigned automation opportunities and improving ETL pipeline performance by 25% while reducing manual handling by 3 hours/day.
- Built REST API integrations and SQL pipelines powering Power BI dashboards, conducting data validation, integration testing, and regression testing to maintain pipeline stability across 6 Agile sprints.

PROJECTS

AI-Powered Credit Card Campaign Analytics System 2025
Python · SQLite · SQL (CTEs, Window Functions) · SciPy · Scikit-learn · Claude API · Streamlit · Tableau Public

- Built end-to-end analytics system on 50,000 synthetic CRM records — Python ETL → SQLite → SQL analyses (CTEs, window functions: channel ROI, spend lift, segment activation) → logistic regression model achieving 88% precision on churn prediction.
- Integrated Claude API narrative engine auto-generating 4-bullet executive summaries with budget recommendations; AI flags 2% spend anomalies (8–15× normal) with plain-English explanations saved to insights.md.
- Deployed Streamlit chatbot for natural-language campaign queries + 5-sheet Tableau Public dashboard (KPI cards, activation funnel, heatmap, spend lift, interactive Story); full SQLite → Python → Claude API → Streamlit + Tableau pipeline demonstrates LLM-on-BI architecture.

GCP Healthcare Data Pipeline – Python Automation & ETL Platform 2023
Google BigQuery · Apache Airflow · Python · SQL · Git · Cloud-Native Architecture · Raw → Staging → Analytics Mart

- Built automated ETL pipeline (Python + BigQuery) with Airflow DAG orchestration ingesting, staging, and loading healthcare appointment data into analytics marts — processes 50K+ records/run at sub-10 min latency.
- Applied layered data quality checks (PK validation, duplicate detection, date sanity) improving reliability by 30%
- Raw → Staging → Analytics Mart architecture supports scalable no-show rate reporting, mirroring enterprise data lakehouse patterns.

Certifications

Databricks Fundamentals , Generative AI Fundamentals , AWS Academy Graduate - Machine Learning Foundations, AWS Academy Graduate - AWS Academy Data Analytics